

How much edit experience should the model receive?

Sequence-edit JEPA validity audit and bounded pilot

17 July 2026

The task is synthetic typo repair

- ▶ Official iGSM supplies a problem and its clean, multi-step worked solution.
- ▶ We damage that solution with token insertions, deletions, and replacements.
- ▶ The expert path is the exact reverse of those corruptions, so it is oracle denoising rather than naturally observed editing.
- ▶ More independent damaged solutions add breadth; more alternative edits per state add local action coverage.

Why it matters: we must establish healthy primitive dynamics before asking whether hierarchy helps.

The old hierarchy round is invalid

- ▶ All five jobs failed before their first optimizer step, so they provide no architectural evidence.
- ▶ The immediate symptom was a 276-token “sentence” exceeding a 128-position table.
- ▶ The cause was deeper: official steps often end in fused tokens such as 6., while the adapter split only on standalone ..
- ▶ The corrected adapter preserves official nested steps and exactly reverses every audited corruption.

What to notice: this is an invalid result, not evidence against hierarchy.

The corrected data contract passes its first gate

Audit check	Observed result
Exact terminal recovery	128/128
Multi-step solution collapse	0/128
Official solution length	median 166.5 tokens
Official steps	median 9
Maximum step length	p95 43 tokens
Repairs per trajectory	mean 11.14
Repair/minimum-distance ratio	mean 1.029

Why it matters: the model now sees the intended step sequence, but these are data-validity observations, not trained-model results.

All six process-valid cells ignore the edit action

Cell	Prediction	No change	Shuffled/matched
512 trajectories	.397	.575	1.00014
2,000 trajectories	.394	.302	1.00000
6,000 trajectories	.224	.113	1.00005
K=1 at 2,000	.287	.268	1.00003
K=4 at 2,000	.175	.123	1.00003

What to notice: lower raw error does not establish dynamics; most cells lose to copying the current state, and shuffling the action changes nothing.

One-token effects disappear in the global target

- ▶ Four exact global next-buffer targets from one state are separated by only .000228 normalized L1.
- ▶ The exact changed-step targets for those edits are separated by .634, about 2,775 times more.
- ▶ The old $K=4$ model assigns changed-step outcomes at 24.7%, chance for four alternatives.
- ▶ The repair predicts exact changed-step anchors for expert and mechanical alternative edits, with no quality or preference labels.
- ▶ We screen coefficients .25, 1, and 4, plus a lower-learning-rate check at the high coefficient.

Next decision: continue only if the correct action beats persistence, shuffled actions, and chance assignment without rank collapse.